

Importing Libraries

In [1]:

```
import random
import pandas as pd
import numpy as np
from scipy.integrate import quad
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
from sklearn.model_selection import cross_val_score, KFold
```

Question 1

In [2]:

```
np.random.seed(1)
x1_train = np.random.gamma(shape=1, scale=1, size=1000)
x2_train = np.random.gamma(shape=1, scale=1, size=1000)
x3_train = np.random.gamma(shape=1, scale=1, size=1000)
w_train = np.random.normal(loc=0, scale=2, size=1000)
y_train = (0.5*x1_train) + (3*x2_train) + (5*x3_train) + (5*x2_train*x3_train) + (2*x1_train*x2_train*x3_train) + w_train

np.random.seed(2)
x1_test = np.random.gamma(shape=1, scale=1, size=1000)
x2_test = np.random.gamma(shape=1, scale=1, size=1000)
x3_test = np.random.gamma(shape=1, scale=1, size=1000)
w_test = np.random.normal(loc=0, scale=2, size=1000)
y_test = (0.5*x1_test) + (3*x2_test) + (5*x3_test) + (5*x2_test*x3_test) + (2*x1_test*x2_test*x3_test) + w_test
```

In [3]:

```
train_df = pd.DataFrame({"x1": x1_train, "x2": x2_train, "x3": x3_train, "y": y_train})
x_train = train_df[["x1", "x2", "x3"]]
y_train = train_df["y"]

test_df = pd.DataFrame({"x1": x1_test, "x2": x2_test, "x3": x3_test, "y": y_test})
x_test = test_df[["x1", "x2", "x3"]]
y_test = test_df["y"]
```

In [4]:

```
lr_model = LinearRegression()
lr_model.fit(x_train, y_train)
y_lr_pred = lr_model.predict(x_test)
lr_mse = mean_squared_error(y_test, y_lr_pred)
print("Linear Regression MSE:", lr_mse)
print("Linear Regression Coefficients:", lr_model.coef_)
```

Linear Regression MSE: 74.56753771040981

Linear Regression Coefficients: [2.19957962 8.38793456 12.19228899]

In [5]:

```
rf_model = RandomForestRegressor(n_estimators = 500)
rf_model.fit(x_train, y_train)
y_rf_pred = rf_model.predict(x_test)
rf_mse = mean_squared_error(y_test, y_rf_pred)
print("Random Forest Regression MSE:", rf_mse)
```

Random Forest Regression MSE: 27.467311656382822

Question 2

In [6]:

```
def function(x):  
    return 3 + (x**2) - (2*np.sin(x))
```

```
a, b = 1, 8  
true_value, error = quad(function, a, b)  
print("True Value:", true_value, u"\u00B1", error)
```

True Value: 189.9617286539798 ± 2.1089988494924473e-12

In [15]:

```
samples = np.random.uniform(low=a, high=b, size=10000)  
function_samples = [function(sample) for sample in samples]
```

```
estimated_value = np.mean(function_samples) * (b-a)  
print("Estimated Value: ", estimated_value)
```

Estimated Value: 190.97185353936558

In [8]:

```
sd = 0  
sd = sum([(result - estimated_value)**2 for result in samples])  
sd = (sd/(10000 - 1))**0.5
```

```
lower = estimated_value - 1.96*sd/(10000**0.5)  
upper = estimated_value + 1.96*sd/(10000**0.5)  
confidence_interval = [lower, upper]  
print("Confidence Interval: ", confidence_interval)
```

Confidence Interval: [184.94448784143083, 192.16113042048005]

Question 6

In [9]:

```
df = pd.DataFrame({"x1": [0, 1, 2, 3, 4], "y": [1, 2, 3, 2, 1]})
```

In [10]:

```
lr_model_1 = LinearRegression()  
lr_model_1.fit(df[["x1"]], df["y"])  
y_pred_1 = lr_model_1.predict(df[["x1"]])  
print("model 1 intercept:", lr_model_1.intercept_)
```

```
lr_model_2 = LinearRegression(fit_intercept=False)  
lr_model_2.fit(df[["x1"]], df["y"])  
y_pred_2 = lr_model_2.predict(df[["x1"]])  
print("model 2 coefficients: ", lr_model_2.coef_)
```

model 1 intercept: 1.8
model 2 coefficients: [0.6]

In [11]:

```
squared_error_1 = ((y_pred_1 - df["y"])**2).mean()  
squared_error_2 = ((y_pred_2 - df["y"])**2).mean()  
print("Average Squared Error for model_1:", squared_error_1, "model_2:", squared_error_2)
```

```
absolute_error_1 = (np.abs(y_pred_1 - df["y"])).mean()  
absolute_error_2 = (np.abs(y_pred_2 - df["y"])).mean()  
print("Average Absolute Error for model_1:", absolute_error_1, "model_2:", absolute_error_2)
```

```
l_15_model_1 = (np.abs(y_pred_1 - df["y"])**1.5).mean()
l_15_model_2 = (np.abs(y_pred_2 - df["y"])**1.5).mean()
print("L1.5 loss for model_1:", l_15_model_1, "model_2:", l_15_model_2)
```

Average Squared Error for model_1: 0.56000000000000002 model_2: 1.64
Average Absolute Error for model_1: 0.64000000000000001 model_2: 1.1600000000000001
L1.5 loss for model_1: 0.5849006163624495 model_2: 1.36348016266711

Question 7

In [12]:

```
#Reading Data
df = pd.read_csv(r"C:\Users\jaske\Desktop\studies\OneDrive\Statistical Methods for Data Science\Assignments\Assignment 1\data.csv")
df["x2"] = df["x2"].astype(int)
df["x2"].unique()
```

Out[12]:

```
array([1, 2, 3])
```

In [13]:

```
#Ordinal Data
ordinal_df = df.copy()
ordinal_model = LinearRegression()
kf = KFold(n_splits=10, shuffle=True, random_state=42)
mse_list = -cross_val_score(ordinal_model, ordinal_df.drop(columns=["y"]), ordinal_df["y"], cv=kf, scoring="neg_mean_squared_error")
mse = (mse_list).mean()
print("Ordered mean squared error:", mse)
```

Ordered mean squared error: 1.2685717076212815

In [14]:

```
#Non Ordinal Data
non_ordinal_df = pd.get_dummies(df, columns=["x2"])
non_ordinal_model = LinearRegression()
kf = KFold(n_splits=10, shuffle=True, random_state=42)
mse_list = -cross_val_score(non_ordinal_model, non_ordinal_df.drop(columns=["y"]), non_ordinal_df["y"], cv=kf, scoring="neg_mean_squared_error")
mse = np.mean(mse_list)
print("UnOrdered mean squared error:", mse)
```

UnOrdered mean squared error: 3.757360926169872e-30

References

Monte Carlo integration in Python. (2022, March 10). GeeksforGeeks. <https://www.geeksforgeeks.org/monte-carlo-integration-in-python/> An Easy Guide to K-Fold Cross-Validation. (2020. November 4). Statology. <https://www.statology.org/k-fold-cross-validation/>